

SKILL DISCOVERY WITH WELL-DEFINED OBJECTIVES (APPENDIX)

Yuu Jinnai, David Abel, Jee Won Park
Department of Computer Science
Brown University
Providence, USA

D Ellis Hershkowitz
Computer Science Department
Carnegie Mellon University
Pittsburgh, USA

Michael Littman, George Konidaris
Department of Computer Science
Brown University
Providence, USA

1 FINDING OPTIONS THAT MINIMIZE PLANNING TIME

Theorem 1.

1. MOMI is $\Omega(\log n)$ hard to approximate even for deterministic MDPs unless $P = NP$.
2. MOMI is $2^{\log^{1-\epsilon} n}$ -hard to approximate for any $\epsilon > 0$ even for deterministic MDP unless $NP \subseteq DTIME(n^{\text{poly} \log n})$.

Proof. First, we show Theorem 4.1 by a reduction from the set cover problem to MOMI with deterministic MDP. We consider two computational problems:

1. MINOPTIONMAXITER (MOMI): Which set of options let value iteration converge in at most ℓ iterations?
2. MINITERMAXOPTION (MIMO): Which set of k or fewer options minimizes the number of iterations to convergence?

More formally, MOMI is defined as follows.

Definition 1 MOMI: *The MINOPTIONMAXITER problem:*

Given an MDP M , a non-negative real-value ϵ , and an integer ℓ , **return** \mathcal{O} that minimizes $|\mathcal{O}|$ subject to $\mathcal{O} \subseteq \mathcal{O}_p$ and $L(\mathcal{O}) \leq \ell$.

We consider a problem OI-DEC which is a decision version of MOMI and MIMO. The problem asks if we can solve the MDP within ℓ iterations using at most k point options.

Definition 2 OI-DEC:

Given an MDP M , a non-negative real-value ϵ , and integers k and ℓ , **return** ‘Yes’ if there exists an option set \mathcal{O} such that $\mathcal{O} \subseteq \mathcal{O}_p$, $|\mathcal{O}| \leq k$ and $L(\mathcal{O}) \leq \ell$. ‘No’ otherwise.

We prove the theorem by reduction from the decision version of the set-cover problem—known to be NP-complete—to OI-DEC. The set-cover problem is defined as follows.

Definition 3 SetCover-DEC:

Given a set of elements \mathcal{U} , a set of subsets $\mathcal{X} = \{X \subseteq \mathcal{U}\}$, and an integer k , **return** ‘Yes’ if there exists a cover $\mathcal{C} \subseteq \mathcal{X}$ that $\bigcup_{X \in \mathcal{C}} X = \mathcal{U}$ and $|\mathcal{C}| \leq k$. ‘No’ otherwise.




Figure 1: Reduction from SetCover-DEC to OI-DEC. The example shows the reduction from an instance of SetCover-DEC which asks if we can pick two subsets from $\mathcal{X} = \{X_1, X_2\}$ where $X_1 = \{1, 2, 3\}$, $X_2 = \{3, 4, 5\}$ to cover all elements $\mathcal{U} = \{1, 2, 3, 4, 5\}$. The SetCover-DEC can be reduced to an instance of OI-DEC where the question is whether the MDP can be solved with 2 iterations of VI by adding at most two point options. The answer of OI-DEC is ‘Yes’ (adding point options from X_1 and X_2 to g will solve the problem), thus the answer of the SetCover-DEC is ‘Yes’. Here the set of initial states corresponds to the cover for the SetCover-DEC.

If there is some $u \in \mathcal{U}$ that is not included in at least one of the subsets X , then the answer is ‘No’. Assuming otherwise, we construct an instance of a shortest path problem (a special case of an MDP problem) as follows (Figure 1). There are four types of states in the MDP: (1) $u_i \in \mathcal{U}$ represents one of the elements in \mathcal{U} , (2) $X_i \in \mathcal{X}$ represents one of the subsets in \mathcal{X} , (3) $X'_i \in \mathcal{X}'$: we make a copy for every state $X_i \in \mathcal{X}$ and call them X'_i , (4) a goal state g . Thus, the state set is $\mathcal{U} \cup \mathcal{X} \cup \mathcal{X}' \cup \{g\}$. We build edges between states as follows: (1) $e(u, X) \in E$ iff $u \in X$: For $u \in \mathcal{U}$ and $X \in \mathcal{X}$, there is an edge between u and X . (2) $\forall X_i \in \mathcal{X}$, $e(X_i, X'_i) \in E$: For every $X_i \in \mathcal{X}$, we have an edge from X_i to X'_i . (3) $\forall e(X', g) \in E$: for every $X' \in \mathcal{X}'$ we have an edge from X' to the goal g . This construction can be done in polynomial time.

Let M be the MDP constructed in this way. We show that $\text{SetCover}(\mathcal{U}, \mathcal{X}, k) = \text{OI-DEC}(M, k, 2)$. Note that by construction every state s_i , s'_i , and g converges to its optimal value within 2 iterations as it reaches the goal state g within 2 steps. A state $u \in \mathcal{U}$ converges within 2 steps if and only if there exists a point option (a) from X to g where $u \in X$, (b) from u to X' where $u \in X$, or (c) from u to g . For options of type (b) and (c), we can find an option of type (a) that makes u converge within 2 steps by setting the initial state of the option to $\mathcal{I}_o = X$, where $u \in X$, and the termination state to $\beta_o = g$. Let \mathcal{O} be the solution of $\text{OI-DEC}(M, k, 2)$. If there exists an option of type (b) or (c), we can swap them with an option of type (a) and still maintain a solution. Let \mathcal{C} be a set of initial states of each option in \mathcal{O} ($\mathcal{C} = \{\mathcal{I}_o | o \in \mathcal{O}\}$). This construction exactly matches the solution of the SetCover-DEC.

□

For Theorems 4.2 and 4.3 we reduce our problem to the Min-Rep, problem, originally defined by ?. Min-Rep is a variant of the better studied label cover problem ? and has been integral to recent hardness of approximation results in network design problems ?. Roughly, Min-Rep asks how to assign as few labels as possible to nodes in a bipartite graph such that every edge is “satisfied.”

Definition 4 Min-Rep:

Given a bipartite graph $G = (A \cup B, E)$ and alphabets Σ_A and Σ_B for the left and right sides of G respectively. Each $e \in E$ has associated with it a set of pairs $\pi_e \subseteq \Sigma_A \times \Sigma_B$ which satisfy it. **Return** a pair of assignments $\gamma_A : A \rightarrow \mathcal{P}(\Sigma_A)$ and $\gamma_B : B \rightarrow \mathcal{P}(\Sigma_B)$ such that for every $e = (A_i, B_j) \in E$ there exists an $(a, b) \in \pi_e$ such that $a \in \gamma_A(A_i)$ and $b \in \gamma_B(B_j)$. The objective is to minimize $\sum_{A_i \in A} |\gamma_A(A_i)| + \sum_{B_j \in B} |\gamma_B(B_j)|$.

We illustrate a feasible solution to an instance of Min-Rep in Figure 2.

The crucial property of Min-Rep we use is that no polynomial-time algorithm can approximate Min-Rep well. Let $\tilde{n} = |A| + |B|$.




Figure 2: An instance of Min-Rep with $\Sigma_A = \{a_1, a_2, a_3\}$ and $\Sigma_B = \{b_1, b_2, b_3\}$. Edge e is labeled with pairs in π_e . Feasible solution (γ_A, γ_B) illustrated where $\gamma_A(A_i)$ and $\gamma_B(B_j)$ below A_i and B_j in blue. Constraints colored to coincide with stochastic action colors in Figure 3.

Lemma 1 (? ?). *Unless $NP \subseteq DTIME(n^{\text{poly} \log n})$, Min-Rep admits no $2^{\log^{1-\epsilon} \tilde{n}}$ polynomial-time approximation algorithm for any $\epsilon > 0$.*

As a technical note, we emphasize that all relevant quantities in Min-Rep are polynomially-bounded. In Min-Rep we have $|\Sigma_A|, |\Sigma_B| \leq \tilde{n}^{c'}$ for constant c' . It immediately follows that $\sum_e |\pi_e| \leq n^c$ for constant c .

1.1 HARDNESS OF APPROXIMATION OF MOMI WITH DETERMINISTIC MDP

Theorem 4.1 Proof. The optimization version of the set-cover problem cannot be approximated within a factor of $c \cdot \ln n$ by a polynomial-time algorithm unless $P = NP$?. The set-cover optimization problem can be reduced to MOMI with a similar construction for a reduction from SetCover-DEC to OI-DEC. Here, the targeted minimization values of the two problems are equal: $P(\mathcal{C}) = |\mathcal{C}|$, and the number of states in OI-DEC is equal to the number of elements in the set cover on transformation. Assume there is a polynomial-time algorithm within a factor of $c \cdot \ln n$ approximation for MOMI where n is the number of states in the MDP. Let SetCover(\mathcal{U}, \mathcal{X}) be an instance of the set-cover problem. We can convert the instance into an instance of MOMI($M, 0, 2$). Using the approximation algorithm, we get a solution \mathcal{O} where $|\mathcal{O}| \leq c \ln n |\mathcal{O}^*|$, where \mathcal{O}^* is the optimal solution. We construct a solution for the set cover \mathcal{C} from the solution to the MOMI \mathcal{O} (see the construction in the proof of Theorem 1). Because $|\mathcal{C}| = |\mathcal{O}|$ and $|\mathcal{C}^*| = |\mathcal{O}^*|$, where \mathcal{C}^* is the optimal solution for the set cover, we get $|\mathcal{C}| = |\mathcal{O}| \leq c \ln n |\mathcal{O}^*| = c \ln n |\mathcal{C}^*|$. Thus, we acquire a $c \cdot \ln n$ approximation solution for the set-cover problem within polynomial time, something only possible if $P=NP$. Thus, there is no polynomial-time algorithm with a factor of $c \cdot \ln n$ approximation for MOMI, unless $P=NP$. \square

1.2 HARDNESS OF APPROXIMATION OF MOMI

We now show our hardness of approximation of $2^{\log^{1-\epsilon} n}$ for MOMI, Theorem 4.2.¹

We start by describing our reduction from an instance of Min-Rep to an instance of MOMI. The intuition behind our reduction is that we can encode choosing a label for a vertex in Min-Rep as choosing an option in our MOMI instance. In particular, we will have a state for each edge in our Min-Rep instance and reward will propagate quickly to that state when value iteration is run only if the options corresponding to a satisfying assignment for that edge are chosen.

More formally, our reduction is as follows. Consider an instance of Min-Rep, MR, given by $G = (A \cup B, E)$, Σ_A, Σ_B and $\{\pi_e\}$. Our instance of MOMI is as follows where $\gamma = 1$ and $l = 3$.²

- **State space** We have a single goal state S_g along with states S'_g and S''_g . For each edge e we create a state S_e . Let $\text{Sat}_A(e)$ consist of all $a \in \Sigma_A$ such that a is in some assignment

¹We assume that \mathcal{O}' is a “good” set of options in the sense that there exists some set $\mathcal{O}^* \subseteq \mathcal{O}'$ such that $L(\mathcal{O}^*) \leq \ell$. We also assume, without loss of generality, that $\epsilon < 1$ throughout this section; other values of ϵ can be handled by re-scaling rewards in our reduction.

²It is easy to generalize these results to $l \geq 4$ by replacing certain edges with paths.




Figure 3: Our MOMI reduction applied to the Min-Rep problem in Figure 2. $e_1 = (A_1, B_1)$, $e_2 = (A_1, B_2)$, $e_3 = (A_2, B_2)$. Actions given in solid lines and each option in \mathcal{O}' represented in its own color as a dashed line from initiation to termination states. Notice that a single option goes from S_{e3b1} and S_{e2b1} to S_g .

in π_e . Define $\text{Sat}_B(e)$ symmetrically. For each edge $e \in E$ we create a set of $2 \cdot |\text{Sat}_A(e)|$ states, namely S_{ea} and S'_{ea} for every $a \in \text{Sat}_A(e)$. We do the same for $b \in \text{Sat}_B(e)$.

- **Actions and Transitions** We have a single action from S'_g to S_g , a single action from S''_g to S'_g . For each edge e we have the following deterministic actions: Every S'_{ea} has a single outgoing action to S_{ea} for $a \in \text{Sat}_A(e)$; Every S_{eb} has a single outgoing action to $S_{eb'}$ for $b \in \text{Sat}_B(e)$; Every S_{ea} has an outgoing action to S_{eb} if $(a, b) \in \pi_e$ and every S'_{eb} has a single outgoing action to S_g ; Lastly, we have a single action from S'_{ea} to S''_g for every $a \in \text{Sat}_A(e)$.
- **Reward** The reward of arriving in S_g is 1. The reward of arriving in every other state is 0.
- **Option Set** Our option set \mathcal{O}' is as follows. For each vertex $A_i \in A$ and each $a \in \Sigma_A$ we have an option $O(A_i, a)$: The initiation set of this option is every S_e where e is incident to A_i ; The termination set of this option is every S_{ea} where A_i is incident to e ; The policy of this option takes the action from S'_{ea} to S_{ea} when in S'_{ea} and the action from S_e to S'_{ea} when in S_e .
Symmetrically, for every vertex $B_j \in B$ and each $b \in \Sigma_B$ we have an option $O(B_j, b)$: The initiation set of this option is every S_{eb} where e is incident to B_j ; The termination set of this option is S_g ; The policy of this option takes the action from S_{eb} to S'_{eb} when in S_{eb} and from S'_{eb} to S_g when in S'_{eb} .

One should think of choosing option $O(v, x)$ as corresponding to choosing label x for vertex v in the input Min-Rep instance. Let $\text{MOMI}(\text{MR})$ be the MDP output given instance MR of Min-Rep and see Figure 3 for an illustration of our reduction.

Let OPT_{MOMI} be the value of the optimal solution to $\text{MOMI}(\text{MR})$ and let OPT_{MR} be the value of the optimal Min-Rep solution to MR. The following lemmas demonstrates the correspondence between a MOMI and Min-Rep solution.

Lemma 2. $\text{OPT}_{\text{MOMI}} \leq \text{OPT}_{\text{MR}}$

Proof. Given a solution (γ_A, γ_B) to MR, define $\mathcal{O}_{\gamma_A, \gamma_B} := \{O(v, x) : v \in V(G) \wedge (\gamma_A(v) = x \vee \gamma_B(v) = x)\}$ as the corresponding set of options. Let γ_A^* and γ_B^* be the optimal solutions to MR which is of cost OPT_{MR} .

We now argue that $\mathcal{O}_{\gamma_A^*, \gamma_B^*}$ is a feasible solution to $MOMI(\text{MR})$ of cost OPT_{MR} , demonstrating that the optimal solution to $MOMI(\text{MR})$ has cost at most OPT_{MR} . To see this notice that by construction the MOMI cost of $\mathcal{O}_{\gamma_A^*, \gamma_B^*}$ is exactly the Min-Rep cost of (γ_A^*, γ_B^*) .

We need only argue, then, that $\mathcal{O}_{\gamma_A^*, \gamma_B^*}$ is feasible for $MOMI(\text{MR})$ and do so now. The value of every state in $MOMI(\text{MR})$ is 1. Thus, we must guarantee that after 3 iterations of value iteration, every state has value 1. However, without any options every state except each S_e has value 1 after 3 iterations of value iteration. Thus, it suffices to argue that $\mathcal{O}_{\gamma_A^*, \gamma_B^*}$ guarantees that every S_e will have value 1 after 3 iterations of value iteration. Since (γ_A^*, γ_B^*) is a feasible solution to MR we know that for every $e = (A_i, B_j)$ there exists an $\bar{a} \in \gamma_A^*(A_i)$ and $\bar{b} \in \gamma_B^*(B_j)$ such that $(\bar{a}, \bar{b}) \in \pi_e$; correspondingly there are options $O(A_i, \bar{a}), O(B_j, \bar{b}) \in \mathcal{O}_{\gamma_A^*, \gamma_B^*}$. It follows that, given options $\mathcal{O}_{\gamma_A^*, \gamma_B^*}$ from S_e one can take option $O(A_i, \bar{a})$ then the action from $S_{e\bar{a}}$ to $S_{e\bar{b}}$ and then option $O(B_j, \bar{b})$ to arrive in S_g ; thus, after 3 iterations of value iteration the value of S_e is 1. Thus, we conclude that after 3 iterations of value iteration every state has converged on its value. \square

We now show that a solution to $MOMI(\text{MR})$ corresponds to a solution to MR. For the remainder of this section $\gamma_A^\mathcal{O}(A_i) := \{a : O(A_i, a) \in \mathcal{O}\}$ and $\gamma_B^\mathcal{O}(B_j) := \{b : O(B_j, b) \in \mathcal{O}\}$ is the Min-Rep solution corresponding to option set \mathcal{O} .

Lemma 3. *For a feasible solution to $MOMI(\text{MR})$, \mathcal{O} , we have $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ is a feasible solution to MR of cost $|\mathcal{O}|$.*

Proof. Notice that by construction the Min-Rep cost of $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ is exactly $|\mathcal{O}|$. Thus, we need only prove that $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ is a feasible solution for MR.

We do so now. Consider an arbitrary edge $e = (A_i, B_j) \in E$; we wish to show that $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ satisfies e . Since \mathcal{O} is a feasible solution to $MOMI(\text{MR})$ we know that after 3 iterations of value iteration every state must converge on its value. Moreover, notice that the value of every state in $MOMI(\text{MR})$ is 1. Thus, it must be the case that for every S_e there exists a path of length 3 from S_e to S_g using either options or actions. The only such paths are those that take an option $O(A_i, a)$, then an action from S_{ea} to S_{eb} then option $O(B_j, b)$ where $(a, b) \in \pi_e$. It follows that $a \in \gamma_A^\mathcal{O}(A_i)$ and $b \in \gamma_B^\mathcal{O}(B_j)$. But since $(a, b) \in \pi_e$, we then know that e is satisfied. Thus, every edge is satisfied and so $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ is a feasible solution to MR. \square

Theorem 4.2 Proof. Assume $\text{NP} \not\subseteq \text{DTIME}(n^{\text{poly log } n})$ and for the sake of contradiction that there exists an $\varepsilon > 0$ for which polynomial-time algorithm \mathcal{A}_{MOMI} can $2^{\log^{1-\varepsilon} n}$ -approximate MOMI. We use \mathcal{A}_{MOMI} to $2^{\log^{1-\varepsilon'} \tilde{n}}$ approximate Min-Rep for a fixed constant $\varepsilon' > 0$ in polynomial-time, thereby contradicting Lemma 1. Again, \tilde{n} is the number of vertices in the graph of the Min-Rep instance.

We begin by noting that the relevant quantities in $MOMI(\text{MR})$ are polynomially-bounded. Notice that the number of states n in the MDP in $MOMI(\text{MR})$ is at most $O(\tilde{n}^2 |\Sigma_A| |\Sigma_B|) = \tilde{n}^c$ for some fixed constant c by the aforementioned assumption that Σ_A and Σ_B are polynomially-bounded in \tilde{n} .³

Our polynomial-time approximation algorithm to approximate instance MR of Min-Rep is as follows: Run \mathcal{A}_{MOMI} on $MOMI(\text{MR})$ to get back option set \mathcal{O} . Return $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ as defined above as our solution to MR.

We first argue that our algorithm is polynomial-time in \tilde{n} . However, notice that for each vertex, we create a polynomial number of states. Thus, the number of states in $MOMI(\text{MR})$ is polynomially-bounded in \tilde{n} and so \mathcal{A}_{MOMI} runs in time polynomial in \tilde{n} . A polynomial runtime of our algorithm immediately follows.

We now argue that our algorithm is a $2^{\log^{1-\varepsilon'} \tilde{n}}$ -approximation for Min-Rep for some $\varepsilon' > 0$. Applying Lemma 3, the approximation of \mathcal{A}_{MOMI} and then Lemma 2, we have that $(\gamma_A^\mathcal{O}, \gamma_B^\mathcal{O})$ is a

³It is also worth noticing that since we create at most $O(\tilde{n}|\Sigma_A| + \tilde{n}|\Sigma_B|)$ options, the total number of options in \mathcal{O}' is at most polynomial in \tilde{n} .

feasible solution for MR with cost

$$\begin{aligned} \text{cost}_{\text{Min-Rep}}(\gamma_A^{\mathcal{O}}, \gamma_B^{\mathcal{O}}) &= |\mathcal{O}| \\ &\leq 2^{\log^{1-\varepsilon} n} \text{OPT}_{\text{MOMI}} \\ &\leq 2^{\log^{1-\varepsilon} n} \text{OPT}_{\text{MR}} \end{aligned}$$

Thus, $(\gamma_A^{\mathcal{O}}, \gamma_B^{\mathcal{O}})$ is a $2^{\log^{1-\varepsilon} n}$ approximation for the optimal Min-Rep solution where n is the number of states in the MDP of $\text{MOMI}(\text{MR})$. Now recalling that $n \leq \tilde{n}^c$ for fixed constant c . We therefore have that $(\gamma_A^{\mathcal{O}}, \gamma_B^{\mathcal{O}})$ is a $2^{\log^{1-\varepsilon} \tilde{n}^c} = 2^{c^{1-\varepsilon} \log^{1-\varepsilon} \tilde{n}} \leq c' \cdot 2^{\log^{1-\varepsilon} \tilde{n}}$ approximation for a constant c' . Choosing ε sufficiently small, we have that $c' \cdot 2^{\log^{1-\varepsilon} \tilde{n}} \leq 2^{\log^{1-\varepsilon'} \tilde{n}}$ for sufficiently large \tilde{n} .

Thus, our polynomial-time algorithm is a $2^{\log^{1-\varepsilon'} \tilde{n}}$ -approximation for Min-Rep for $\varepsilon' > 0$, thereby contradicting Lemma 1. We conclude that MOMI cannot be $2^{\log^{1-\varepsilon} n}$ -approximated. \square

Theorem 2. A-MOMI has the following properties:

1. A-MOMI runs in polynomial time.
2. It guarantees that the MDP is solved within ℓ iterations using the option set acquired by A-MOMI \mathcal{O} .
3. If the MDP is deterministic, the option set is at most $\max_{s \in \mathcal{S}} X_s$ times larger than the smallest option set possible to solve the MDP within ℓ iterations.

Theorem 2.1. A-MOMI runs in polynomial time.

Proof. Each step of the procedure runs in polynomial time.

- (1) Solving an MDP takes polynomial time ?. To compute d we need to solve MDPs at most $|\mathcal{S}|$ times. Thus, it runs in polynomial time.
- (4) We solve the set cover using a polynomial time approximation algorithm ? which runs in $O(|\mathcal{S}|^3)$, thus run in polynomial time.
- (2), (3), and (5) Immediate. \square

Theorem 2.2. A-MOMI guarantees that the MDP is solved within ℓ iterations using the option set \mathcal{O} .

Proof. A state $s \in X_g^+$ reaches optimal within ℓ steps by definition. For every state $s \in \mathcal{S} \setminus X_g^+$, the set cover guarantees that we have $X_{s'} \in \mathcal{C}$ such that $d(s, s') < \ell$. As we generate an option from s' to g , s' reaches to optimal value with 1 step. Thus, s reaches to ε -optimal value within $d(s, s') + 1 \leq \ell$. Therefore, every state reaches ε -optimal value within ℓ steps. \square

Theorem 2.3. If the MDP is deterministic, the option set is at most $\max_{s \in \mathcal{S}} X_s$ times larger than the smallest option set possible to solve the MDP within ℓ iterations.

Proof. Using a suboptimal algorithm by ?? we get \mathcal{C} such that $|\mathcal{C}| \leq \Delta |\mathcal{C}|$ where Δ is the maximum size of subsets in \mathcal{X} . Thus, $|\mathcal{O}| = |\mathcal{C}| \leq \Delta |\mathcal{C}^*| = \Delta |\mathcal{O}^*|$. \square

1.3 A-MIMO

The approximation algorithm for MIMO (A-MIMO) is as follows.

1. Compute an asymmetric distance function $d_\varepsilon(s, s') : \mathcal{S} \times \mathcal{S} \rightarrow \mathbb{N}$ representing the number of iterations for a state s to reach its ε -optimal value if we add a point option from a state s' to a goal state g .
2. Using this distance function, solve an asymmetric k -center problem, which finds a set of center states that minimizes the maximum number of iterations for every state to converge.

3. Generate point options with initiation states set to the center states in the solution of the asymmetric k -center, and termination states set to the goal.

Theorem 2.4. *A-MIMO runs in polynomial time.*

Proof. Each step of the procedure runs in polynomial time.

(1) Solving an MDP takes polynomial time. To compute d we need to solve MDPs at most $|S|$ times. Thus, it runs in polynomial time.

(2) The approximation algorithm we deploy for solving the asymmetric- k center which runs in polynomial time. Because the procedure by ? ? terminates immediately after finding a set of options which guarantees the suboptimality bounds, it tends to find a set of options smaller than k . In order to use the rest of the options effectively within polynomial time, we use a procedure Expand to greedily add a few options at once until it finds all k options. We enumerate all possible set of options of size $r = \lceil \log k \rceil$ (if $|\mathcal{O}| + \log k > k$ then we set $r = k - |\mathcal{O}|$) and add a set of options which minimizes ℓ (breaking ties randomly) to the option set \mathcal{O} . We repeat this procedure until $|\mathcal{O}| = k$. This procedure runs in polynomial time. The number of possible option set of size r is ${}^r C_n = O(n^r) = O(k)$. We repeat this procedure at most $\lceil k / \log k \rceil$ times, thus the total computation time is bounded by $O(k^2 / \log k)$.

(3) Immediate.

Therefore, A-MIMO runs in polynomial time. \square

Before we show that it is sufficient to consider a set of options with its terminal state set to the goal state of the MDP.

Lemma 4. *There exists an optimal option set for MIMO and MOMI with all terminal state set to the goal state.*

Proof. Assume there exists an option with terminal state set to a state other than the goal state in the optimal option set \mathcal{O} . By triangle inequality, swapping the terminal state to the goal state will monotonically decrease $d(s, g)$ for every state. By swapping every such option we can construct an option set \mathcal{O}' with $L(\mathcal{O}') \leq L(\mathcal{O})$. \square

Lemma imply that discovering the best option set among option sets with their terminal state fixed to the goal state is sufficient to find the best option set in general. Therefore, our algorithms seek to discover options with termination state fixed to the goal state.

Using the option set acquired, the number of iterations to solve the MDP is bounded by $P(\mathcal{C})$. To prove this we first generalize the definition of the distance function to take a state and a set of states as arguments $d_\epsilon : \mathcal{S} \times 2^{\mathcal{S}} \rightarrow \mathbb{N}$. Let $d_\epsilon(s, \mathcal{C})$ the number of iterations for s to converge ϵ -optimal if every state $s' \in \mathcal{C}$ has converged to ϵ -optimal: $d_\epsilon(s, \mathcal{C}) := \min(d'_\epsilon(s), 1 + d'_\epsilon(s, \mathcal{C})) - 1$. As adding an option will never make the number of iterations larger,

Lemma 5.

$$d(s, \mathcal{C}) \leq \min_{s' \in \mathcal{C}} d(s, s'). \quad (1)$$

Using this, we show the following proposition.

Theorem 2.5. *The number of iterations to solve the MDP using the acquired options is upper bounded by $P(\mathcal{C})$.*

Proof. $P(\mathcal{C}) = \max_{s \in \mathcal{S}} \min_{c \in \mathcal{C}} d(s, c) \geq \max_{s \in \mathcal{S}} d(s, \mathcal{C}) = L(\mathcal{O})$ (using Equation 2). Thus $P(\mathcal{C})$ is an upper bound for $L(\mathcal{O})$. \square

The reason why $P(\mathcal{C})$ does not always give us the exact number of iterations is because adding two options starting from s_1, s_2 may make the convergence of s_0 faster than $d(s_0, s_1)$ or $d(s_0, s_2)$. Example: Figure 5 is an example of such an MDP. From s_0 it may transit to s_1 and s_2 with probability 0.5 each. Without any options, the value function converges to exactly optimal value for every state with 3 steps. Adding an option either from s_1 or s_2 to g does not shorten the iteration for s_0 to




Figure 4: An example of an MDP where $d(s, \mathcal{C}) < \min_{s' \in \mathcal{C}} d(s, s')$. Here the transition induced by the optimal policy is stochastic, thus from s_0 one may go to s_1 and s_2 by probability 0.5 each. Either adding an option from s_1 or s_2 to g does not make the convergence faster, but adding both makes it faster.

converge. However, if we add two options from s_1 and s_2 to g , s_0 converges within 2 steps, thus the MDP is solved with 2 steps.

The equality of the statement 2 holds if the MDP is deterministic. That is, $d(s, \mathcal{C}) = \min_{s' \in \mathcal{C}} d(s, s')$ for deterministic MDP.

Theorem 2.6.

If the MDP is deterministic, it has a bounded suboptimality of $\log^ k$.*

Proof. First we show $P(\mathcal{C}^*) = L(\mathcal{O}^*)$ for deterministic MDP. From $d(s, \mathcal{C}) = \min_{s' \in \mathcal{C}} d(s, s')$, $P(\mathcal{C}^*) = \max_{s \in \mathcal{S}} \min_{c \in \mathcal{C}^*} d(s, c) = \max_{s \in \mathcal{S}} d(s, \mathcal{C}^*) = L(\mathcal{O}^*)$.

The asymmetric k -center solver guarantees that the output \mathcal{C} satisfies $P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*)$ where n is the number of nodes ?. Let $\text{MIMO}(M, \epsilon, k)$ be an instance of MIMO. We convert this instance to an instance of asymmetric k -center $\text{AsymKCenter}(\mathcal{U}, d, k)$, where $|\mathcal{U}| = |\mathcal{S}|$. By solving the asymmetric k -center with the approximation algorithm, we get a solution \mathcal{C} which satisfies $P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*)$. Thus, the output of the algorithm \mathcal{O} satisfies $L(\mathcal{O}) = P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*) = c(\log^* k + O(1))L(\mathcal{O}^*)$. Thus, $L(\mathcal{O}) \leq c(\log^* k + O(1))L(\mathcal{O}^*)$ is derived. \square

Before we show that it is sufficient to consider a set of options with its terminal state set to the goal state of the MDP.

Lemma 6. *There exists an optimal option set for MIMO and MOMI with all terminal state set to the goal state.*

Proof. Assume there exists an option with terminal state set to a state other than the goal state in the optimal option set \mathcal{O} . By triangle inequality, swapping the terminal state to the goal state will monotonically decrease $d(s, g)$ for every state. By swapping every such option we can construct an option set \mathcal{O}' with $L(\mathcal{O}') \leq L(\mathcal{O})$. \square

Lemma imply that discovering the best option set among option sets with their terminal state fixed to the goal state is sufficient to find the best option set in general. Therefore, our algorithms seek to discover options with termination state fixed to the goal state.

Using the option set acquired, the number of iterations to solve the MDP is bounded by $P(\mathcal{C})$. To prove this we first generalize the definition of the distance function to take a state and a set of states as arguments $d_\epsilon : \mathcal{S} \times 2^{\mathcal{S}} \rightarrow \mathbb{N}$. Let $d_\epsilon(s, \mathcal{C})$ the number of iterations for s to converge ϵ -optimal if every state $s' \in \mathcal{C}$ has converged to ϵ -optimal: $d_\epsilon(s, \mathcal{C}) := \min(d'_\epsilon(s), 1 + d'_\epsilon(s, \mathcal{C})) - 1$. As adding an option will never make the number of iterations larger,

Lemma 7.

$$d(s, \mathcal{C}) \leq \min_{s' \in \mathcal{C}} d(s, s'). \quad (2)$$




Figure 5: An example of an MDP where $d(s, \mathcal{C}) < \min_{s' \in \mathcal{C}} d(s, s')$. Here the transition induced by the optimal policy is stochastic, thus from s_0 one may go to s_1 and s_2 by probability 0.5 each. Either adding an option from s_1 or s_2 to g does not make the convergence faster, but adding both makes it faster.

Using this, we show the following proposition.

Theorem 2.7. *The number of iterations to solve the MDP using the acquired options is upper bounded by $P(\mathcal{C})$.*

Proof. $P(\mathcal{C}) = \max_{s \in \mathcal{S}} \min_{c \in \mathcal{C}} d(s, c) \geq \max_{s \in \mathcal{S}} d(s, \mathcal{C}) = L(\mathcal{O})$ (using Equation 2). Thus $P(\mathcal{C})$ is an upper bound for $L(\mathcal{O})$. \square

The reason why $P(\mathcal{C})$ does not always give us the exact number of iterations is because adding two options starting from s_1, s_2 may make the convergence of s_0 faster than $d(s_0, s_1)$ or $d(s_0, s_2)$. Example: Figure 5 is an example of such an MDP. From s_0 it may transit to s_1 and s_2 with probability 0.5 each. Without any options, the value function converges to exactly optimal value for every state with 3 steps. Adding an option either from s_1 or s_2 to g does not shorten the iteration for s_0 to converge. However, if we add two options from s_1 and s_2 to g , s_0 converges within 2 steps, thus the MDP is solved with 2 steps.

The equality of the statement 2 holds if the MDP is deterministic. That is, $d(s, \mathcal{C}) = \min_{s' \in \mathcal{C}} d(s, s')$ for deterministic MDP.

Theorem 2.8.

If the MDP is deterministic, it has a bounded suboptimality of $\log^ k$.*

Proof. First we show $P(\mathcal{C}^*) = L(\mathcal{O}^*)$ for deterministic MDP. From $d(s, \mathcal{C}) = \min_{s' \in \mathcal{C}} d(s, s')$, $P(\mathcal{C}^*) = \max_{s \in \mathcal{S}} \min_{c \in \mathcal{C}^*} d(s, c) = \max_{s \in \mathcal{S}} d(s, \mathcal{C}^*) = L(\mathcal{O}^*)$.

The asymmetric k -center solver guarantees that the output \mathcal{C} satisfies $P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*)$ where n is the number of nodes. Let $\text{MIMO}(M, \epsilon, k)$ be an instance of MIMO. We convert this instance to an instance of asymmetric k -center $\text{AsymKCenter}(\mathcal{U}, d, k)$, where $|\mathcal{U}| = |\mathcal{S}|$. By solving the asymmetric k -center with the approximation algorithm, we get a solution \mathcal{C} which satisfies $P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*)$. Thus, the output of the algorithm \mathcal{O} satisfies $L(\mathcal{O}) = P(\mathcal{C}) \leq c(\log^* k + O(1))P(\mathcal{C}^*) = c(\log^* k + O(1))L(\mathcal{O}^*)$. Thus, $L(\mathcal{O}) \leq c(\log^* k + O(1))L(\mathcal{O}^*)$ is derived. \square

2 FINDING OPTIONS THAT MINIMIZE LEARNING TIME FOR HARD-EXPLORATION TASKS

Theorem 3. *Assume a stochastic shortest path problem to reach a goal g where a non-positive reward $r_c \leq 0$ is given for non-goal states and $\gamma = 1$. Let P be a random walk transition matrix: $P(s, s') = \sum_{a \in A} \pi(s) T(s, a, s')$:*

$$\forall g : V_g^\pi(s) \geq r_c \mathbb{E}[C(G)],$$

where $C(G) = \max_{s \in \mathcal{S}} C_s(G)$ and $C_s(G)$ is a cover time of a transition matrix P starting from state s .

Proof. The value of state s is r_c times the expected number of steps to reach the goal state. Thus,

$$\begin{aligned} V_g^\pi(s) &= r_c \mathbb{E}[H_{sg}] \\ &\geq r_c \mathbb{E}[\max_{s' \in S} H_{ss'}] \\ &= r_c \mathbb{E}[C_s(G)] \\ &\geq r_c \mathbb{E}[C(G)] \end{aligned}$$

□

Theorem 4. Assume that a random walk induced by a policy π is a uniform random walk. Adding two options by the algorithm improves the upper bound of the cover time if the multiplicity of the second smallest eigenvalue is one:

$$\mathbb{E}[C(G')] \leq \frac{n^2 \ln n}{\lambda_2(\mathcal{L}) + F} (1 + o(1)), \quad (3)$$

where $\mathbb{E}[C(G')]$ is the expected cover time of the augmented graph, $F = \frac{(v_i - v_j)^2}{6/(\lambda_3 - \lambda_2 + 3/2)}$, and v_i, v_j are the maximum and minimum values of the Fiedler vector. If the multiplicity of the second smallest eigenvalue is more than one, then adding any single option cannot improve the algebraic connectivity.

Proof. Assume the multiplicity of the second smallest eigenvalue is one. Let \mathcal{L}' be the graph Laplacian of the graph with an edge inserted to \mathcal{L} using the algorithm by ?? . By adding a single edge, the algebraic connectivity is guaranteed to increase at least by F :

$$\lambda_2 \geq \lambda_2 + \frac{(v_i - v_j)^2}{6/(\lambda_3 - \lambda_2) + 3/2}, \quad (4)$$

and the upper bound of the cover time is guaranteed to decrease:

$$\begin{aligned} \mathbb{E}[C(G')] &\leq \frac{n^2 \ln n}{\lambda_2} (1 + o(1)) \\ &\leq \frac{n^2 \ln n}{\lambda_2 + \frac{(v_i - v_j)^2}{6/(\lambda_3 - \lambda_2) + 3/2}} (1 + o(1)). \end{aligned}$$

As $\frac{(v_i - v_j)^2}{6/(\lambda_3 - \lambda_2) + 3/2}$ is positive,

$$\frac{n^2 \ln n}{\lambda_2 + \frac{(v_i - v_j)^2}{6/(\lambda_3 - \lambda_2) + 3/2}} (1 + o(1)) < \frac{n^2 \ln n}{\lambda_2} (1 + o(1)), \quad (5)$$

thus the upper bound is guaranteed to decrease.

Assume the second smallest eigenvalue is more than one. Then, $\lambda_2(\mathcal{L}) = \lambda_3(\mathcal{L})$. From eigenvalue interlacing ?, for any edge insertion, $\lambda_2(\mathcal{L}) \leq \lambda_2(\mathcal{L}') \leq \lambda_3(\mathcal{L})$. Thus, $\lambda_2(\mathcal{L}') = \lambda_2(\mathcal{L})$. □